



THE LONDON SCHOOL  
OF ECONOMICS AND  
POLITICAL SCIENCE ■

# Week 9: Text as Data

MY472: Data for Data Scientists

<https://lse-my472.github.io/>

Autumn Term 2025

**Ryan Hübert**

Associate Professor of Methodology

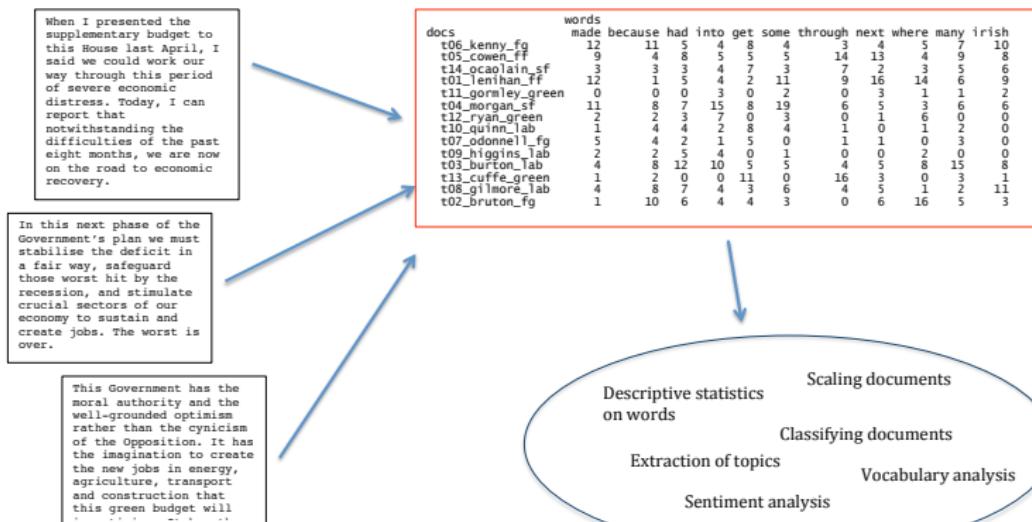
# Outline

- 1 Text as (quantitative) data
- 2 More complex document features
- 3 Getting text data from other formats
- 4 Data security

# Text as (quantitative) data

**Quantitative text analysis (QTA):** use quantitative methods to analyse and work with texts

But first, we need to *quantify* the texts!



# Quantifying texts

We'll be quantifying texts using the `{quanteda}` package in R

- ▷ `{quanteda}` was developed by a former LSE prof, Ken Benoit
- ▷ Offers a seamless package for doing a lot of QTA tasks.

```
install.packages("quanteda")
install.packages("quanteda.textplots")
install.packages("quanteda.textstats")
```

There are other tools in python, probably more common in industry

## Document selection

A **document** is the fundamental unit of analysis for quantitative text analysis, for example:

- ▷  $n$ -word sequences
- ▷ Sentences
- ▷ Pages
- ▷ Paragraphs
- ▷ Natural units (a speech, a poem, a manifesto)
- ▷ Aggregation of units (e.g. all speeches by party and year)

A collection of documents is called a **corpus**

How you define documents depends on your goal

## Side note on terminology

Let's keep some words straight for this lecture:

- ▷ a **document** is the chosen unit of analysis; could be a full “document” in colloquial sense, or not
- ▷ a **text** (used as a countable noun) is what we often refer to as a “document” in the colloquial sense, e.g. a novel, a speech, a legal opinion, a tweet, etc.
- ▷ **text** (used as an uncountable noun) is a general word used to label a type of data, similar to “string”

I'll try to be consistent, but context will usually be informative

## Example: US President Trump's tweets

To demonstrate quantifying texts, we'll use a corpus of US President Trump's tweets

- ▷ The corpus covers 2017 and half of 2018
- ▷ Available on the course website

Data is stored in a `.json` file based on the Twitter API

- ▷ Convenient `{streamR}` package to parse JSON in the Twitter API format

## Example: US President Trump's tweets

Donald J. Trump   
@realDonaldTrump

TO ALL AMERICANS-  
#HappyNewYear & many blessings to you all! Looking forward to a  
wonderful & prosperous 2017 as we work together to #MAGA



**HAPPY NEW YEAR!  
2017**  
.....  
**TO ALL AMERICANS!**  
**HAPPY NEW YEAR & MANY  
BLESSINGS TO YOU ALL!  
LOOKING FORWARD TO A  
WONDERFUL AND  
PROSPEROUS 2017 AS WE  
WORK TOGETHER TO MAKE  
AMERICA GREAT AGAIN!**  
 President-Elect  
Donald J. Trump

5:00 AM · Jan 1, 2017

Source: <https://x.com/realDonaldTrump/status/815422340540547073>

## Example: US President Trump's tweets

```
{ "created_at": ["Sun Jan 01 05:00:10 +0000 2017"],  
  "id": [8.15422340540547e+17],  
  "id_str": ["815422340540547073"],  
  "full_text": ["TO ALL AMERICANS-\n#HappyNewYear & many blessings to  
   you all! Looking forward to a wonderful & prosperous  
   2017 as we work together to #MAGA\u0001F1FA\u0001F1F8  
   https://t.co/UaBFaoDYHe"],  
  "truncated": [false],  
  "display_text_range": [[0],[148]],  
  "entities": {"hashtags": [{"text": "HappyNewYear", "indices": [[18],[31]]},  
    {"text": "MAGA", "indices": [[141],[146]]}],  
    "symbols": [],  
    "user_mentions":[],  
    "urls":[],  
    "media": [{"id": [8.15422333510816e+17],  
      "id_str": ["815422333510815746"],  
      "indices": [[149],[172]],  
      "media_url": ["http://pbs.twimg.com/media/C1D2SsLVE"],  
      "media_url_https": ["https://pbs.twimg.com/media/C1D2SsLVE"],  
      "url": ["https://t.co/UaBFaoDYHe"],  
      ...  
    }]}]
```

# Defining document features

Documents contain **features**, which can be:

- ▷ characters
- ▷ words
- ▷ word “stems” or “lemmas” (more later)
- ▷ word segments, especially for languages using compound words, such as German, e.g. *Saunauntensitzer*
- ▷ “word” sequences, especially when inter-word delimiters (usually white space) are not commonly used, e.g., in Chinese
- ▷ linguistic features, such as parts of speech
- ▷ coded or annotated text segments
- ▷ word embeddings

## The most common approach: bag of words

Most common approach to quantifying text: **bag of words** model

- ▷ Single *words* are the relevant *features* of each document
- ▷ Documents are quantified by counting occurrences of words
- ▷ Word order does not matter
- ▷ Discards grammar and syntax

## Why bag of words?

Most obvious reason: it's very simple

But also, context is often uninformative and conditional on presence of words

- ▷ Individual word usage tends to be associated with a particular degree of affect, position, etc. without regard to context
- ▷ So presence of words by itself captures info about context

Plus, *single* words tend to be the most informative since co-occurrences of multiple words (“*n*-grams”) are relatively rare

In social science applications: bag of words works well most of the time (i.e., it's been validated *a lot*)

# Why bag of words?

There are times where word order is important, e.g.

- ▷ **Text reuse**: plagiarism detection software used for social science applications, such as [Corley \(2007\)](#) and [Grimmer \(2010\)](#)
- ▷ **Parts of speech tagging**: tagging words in documents with grammatical information, with applications such as [Bamman and Smith \(2014\)](#) and [Handler et al \(2016\)](#)
- ▷ **Named entity recognition (NER)**: finding and tagging words in documents that are people, organisations, or places, with applications such as [Copus](#), [Hübert and Pellaton \(2024\)](#)

As usual: whether word order matters depends on the QTA task

# Implementing bag of words

Goal: get from a set of texts to a quantitative dataset

There is a basic four step process for implementing bag of words to quantify texts:

1. Choose unit of analysis
2. Tokenise
3. Reduce complexity
4. Create **document feature matrix**

People often refer to this as **preprocessing**

The “reducing complexity” part is where most of the discretion is

You should justify the preprocessing steps you take

# Tokenising

You start by **tokenising** each document:

- ▷ Split into an array of words, each called a **token**
- ▷ For English (and many other languages), use white space; trickier for logographic languages (e.g. Chinese)
- ▷ Tokens are *mostly* “words”—but not always

A list of all the distinct tokens used in an entire corpus is a **vocabulary**

- ▷ Each element of the vocabulary is a **type**

## Tokenising Trump's tweet

Original (plain) text of Trump tweet:

TO ALL AMERICANS-\n#HappyNewYear & many blessings  
to you all! Looking forward to a wonderful &  
prosperous 2017 as we work together to  
#MAGA\U0001F1FA\U0001F1F8 <https://t.co/UaBFaoDYHe>

## Tokenising Trump's tweet

Use white space to tokenise:

```
c("TO", "ALL", "AMERICANS-", "#HappyNewYear", "&#038;",  
  "many", "blessings", "to", "you", "all!", "Looking",  
  "forward", "to", "a", "wonderful", "&#038;",  
  "prosperous", "2017", "as", "we", "work", "together",  
  "to", "#MAGA", "https://t.co/UaBFaoDYHe")
```

Notice some issues here:

- ▷ some useless punctuation: “AMERICANS-”
- ▷ some “non-words” included: links, ampersands
- ▷ the words “to” and “TO” are considered different words

This will create a vocabulary that is too large and redundant

## Reduce complexity: cleaning up formatting

To deal with these problems, we can: remove “non-words” and punctuation, and make text lowercase

```
c("to", "all", "americans", "#happynewyear", "many",
  "blessings", "to", "you", "all", "looking", "forward",
  "to", "a", "wonderful", "prosperous", "as", "we",
  "work", "together", "to", "#maga")
```

Note:

- ▷ Had to manually get rid of symbols written in [HTML syntax](#) (e.g. &amp;) using regex pattern "[&][#]?[A-z]+;"
- ▷ Decided to keep Twitter hashtags intact (why?)

## Reduce complexity: removing stop words

**Stop words** are words that occur very frequently in a language but do not provide much information

Some very common English stop words are:

a, able, about, across, after, all, almost, also, am, among, an, and, any, are, as, at, be, because, been, but, by, can, cannot, could, dear, did, do, does, either, else, ever, every, for, from, get, got, had, has, have, he, her, hers, him, his, how, however, I, if, in, into, is, it, its, just, least, let, like, likely, may, me, might, most, must, my, neither, no, nor, not, of, off, often, on, only, or, other, our, own, rather, said, say, says, she, should, since, so, some, than, that, the, their, them, then, there, these, they, this, tis, to, too, twas, us, wants, was, we, were, what, when, where, which, while, who, whom, why, will, with, would, yet, you, your

In QTA, it is very common to remove stop words

- ▷ But this depends on your task, see [Pennebaker \(2011\)](#)
- ▷ There are different lists out there, some longer, some shorter

## Removing stop words from Trump's tweet

We can remove all the stop words (defined by the list above) from the Trump tweet:

```
c("americans", "#happynewyear", "many", "blessings",
  "looking", "forward", "wonderful", "prosperous",
  "work", "together", "#maga")
```

## Reduce complexity: stemming and lemmatisation

Many words have multiple “forms” with different spellings, e.g.

- ▷ Tenses: “I see” and “I saw”
- ▷ Pluralisation: “family” and “families”
- ▷ Contractions: “families” and “family’s”

So far, we have treated all these as different words

But, you may wish to create **equivalence classes** of words considered to convey same meaning

- ▷ Basic idea: for every token, identify the “root” word, and replace the token with root word
- ▷ This reduces vocabulary, and can be quite useful for comparing across documents

You may or may not want to do this depending on your task

# Reduce complexity: stemming and lemmatisation

Two common approaches:

1. **Lemmatisation**: refers to the algorithmic process of converting words to their lemma
  - ▷ A word's **lemma** is its “canonical form”
2. **Stemming**: removing the ends of words using a set of rules

Basic difference: stemmers operate on single words without knowledge of the context, lemmatisers are more powerful

- ▷ There is a computational tradeoff

Example: **production**, **producer**, **produce**, **produces**,  
**produced** all replaced with **produc**

Tools for each available in `{quanteda}` package

## Reduce complexity: stemming and lemmatisation

Lots of different ways to stem and lemmatise

**Porter stemmer** is most common, but gets many stems wrong:

- ▷ `policy` and `police` considered (wrongly) equivalent
- ▷ `general` becomes `gener`, `iteration` becomes `iter`

There are other corpus-based, statistical, and mixed approaches designed to overcome these limitations

Plus, sometimes stemming isn't appropriate: [Schofield and Mimno \(2016\)](#) find that “stemmers produce no meaningful improvement in likelihood and coherence (of topic models) and in fact can degrade topic stability”

Take away: you have to read and validate!

## Reduce complexity: stemming and lemmatisation

We can use `{quanteda}`'s default stemming function `tokens_wordstem()` on the Trump tweet, yielding:

```
c("american", "#happynewyear", "mani", "bless", "look",
  "forward", "wonder", "prosper", "work", "togeth",
  "#maga")
```

Note:

- ▷ `{quanteda}` uses the **Snowball stemmer** by default
- ▷ The stemmer did many transformations, e.g.:
  - `americans` → `american`
  - `many` → `mani`
  - `blessing` → `bless`
  - `together` → `togeth`

## The document-feature matrix (DFM)

So far, we've just “preprocessed” one example document

You repeat this process for every document

- ▷ This yields a vocabulary of types for the corpus
- ▷ These are the *features* to be analysed (again: remember we're assuming bag of words!)

Each document uses some of the features in the vocabulary

- ▷ You can count how many times

A **document-feature matrix** (math:  $\mathbf{W}$ ) is a matrix of  $N$  documents (rows) by  $J$  features (columns) where:

- ▷ each  $W_{ij}$  counts the number of times the  $j$ th feature appears in the  $i$ th document

## The document-feature matrix (DFM)

All of Trump's tweets from January 2017:

```
Document-feature matrix of: 212 documents, 1,039 features (98.94%
sparse) and 0 docvars.
```

docs		american	#happynewyear	mani	bless
2017-01-01 05:00:10		1		1	1
2017-01-01 05:39:13		0		1	0
2017-01-01 05:43:23		0		0	1
2017-01-01 05:44:17		0		0	0
2017-01-01 06:49:33		0		0	0

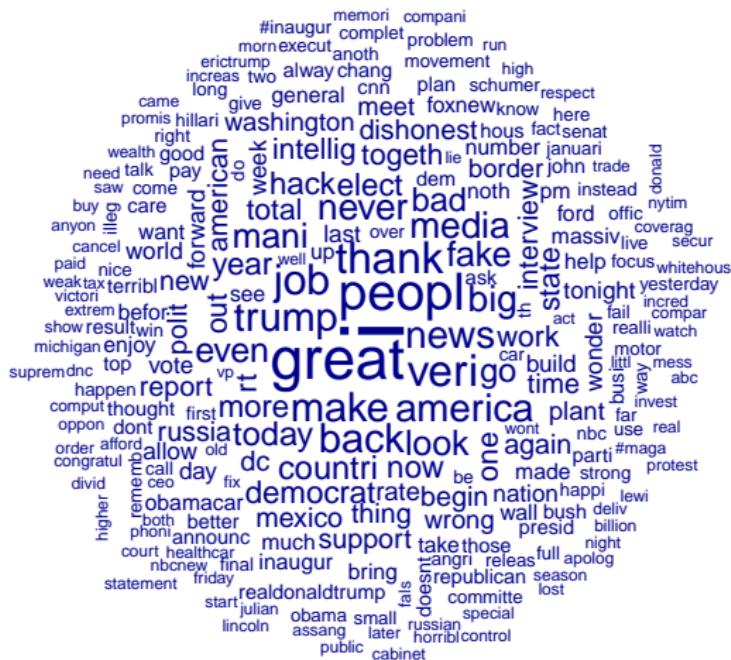
```
[ reached max_ndoc ... 207 more documents, reached max_nfeat ...
1,035 more features ]
```

This DFM has  $J = 1,039$  features and  $N = 212$  documents

- ▷ Each cell is a count, e.g.  $W_{11} = 1$  and  $W_{21} = 0$

# Wordclouds

Wordclouds are basic visualisations of the data in a DFM



# Outline

- 1 Text as (quantitative) data
- 2 More complex document features
- 3 Getting text data from other formats
- 4 Data security

## Not all tokens *should* be unigrams

When we tokenise using white spaces, we create **unigrams**

We could define tokens differently by choosing a **collocation**:

- ▷ **bigrams**: pairs of adjacent words
- ▷ **trigrams**: triples of adjacent words
- ▷ more generally: ***n*-grams**: *n* adjacent words

Example: **capital gains tax** can be represented as

- ▷ unigrams: `c("capital", "gains", "tax")`
- ▷ bigrams: `c("capital gains", "gains tax")`
- ▷ trigrams: `c("capital gains tax")`

## Not all tokens *should* be unigrams

Why would you want to depart from unigrams?

1. You might want to do your entire analysis using  $n$ -grams instead of unigrams
  - ▷ It could be useful for situations where word order matters like simple text completion
2. Many collocations have independent meaning that you might want to retain in your analysis
  - ▷ For example, **United Kingdom** makes more sense left as a bigram than as two unigrams
  - ▷ In cases like this: need to manually define which phrases to leave as  $n$ -grams

How do you decide which words to collocate into  $n$ -grams?

# Important collocations

$C(w^1 w^2)$	$w^1$	$w^2$
80871	of	the
58841	in	the
26430	to	the
21842	on	the
21839	for	the
18568	and	the
16121	that	the
15630	at	the
15494	to	be
13899	in	a
13689	of	a
13361	by	the
13183	with	the
12622	from	the
11428	New	York
10007	he	said
9775	as	a
9231	is	a
8753	has	been
8573	for	a

**Table 5.1** Finding Collocations: Raw Frequency.  $C(\cdot)$  is the frequency of something in the corpus.

# Important collocations

$C(w^1 w^2)$	$w^1$	$w^2$
80871	of	the
58841	in	the
26430	to	the
21842	on	the
21839	for	the
18568	and	the
16121	that	the
15630	at	the
15494	to	be
13899	in	a
13689	of	a
13361	by	the
13183	with	the
12622	from	the
11428	New	York
10007	he	said
9775	as	a
9231	is	a
8753	has	been
8573	for	a

**Table 5.1** Finding Collocations: Raw Frequency.  $C(\cdot)$  is the frequency of something in the corpus.

## Identifying collocations in {quanteda}

Find frequent collocations with: `textstat_collocations()`

Let's find the common bigrams in the Trump tweet corpus:

```
# A tibble: 5,335 x 6
  collocation    count count_nested length lambda      z
  <chr>        <int>        <int>   <dbl>   <dbl>   <dbl>
1 fake news     217          0       2     8.01   40.1
2 tax cut       130          0       2     7.11   35.8
3 make america  99           0       2     5.60   35.6
4 unit state    101          0       2     7.25   30.6
5 north korea   116          0       2     9.32   30.0
6 news media    63           0       2     5.02   28.7
7 america great 91           0       2     4.00   28.5
8 great again   98           0       2     4.75   28.2
9 illeg immigr  39           0       2     6.56   26.1
10 work hard    46           0       2     5.42   25.9
# i 5,325 more rows
```

## Word embeddings

Bag of words: each word has a distinct, unique meaning

- ▷ In math terms: we treat words as **one-hot encodings**

Concrete example: consider a vocabulary of three words:

`c("cat", "dog", "rat")`

Can represent each word in **vector format**:

- ▷ The word `cat` is  $(1, 0, 0)$
- ▷ The word `dog` is  $(0, 1, 0)$
- ▷ The word `rat` is  $(0, 0, 1)$

(Should be obvious why this is called “one-hot encoding”)

These are *orthogonal*: zero similarity between the words

## Word embeddings

But what if words are actually representations of a smaller group of concepts, where each word is a *mix* of concepts?

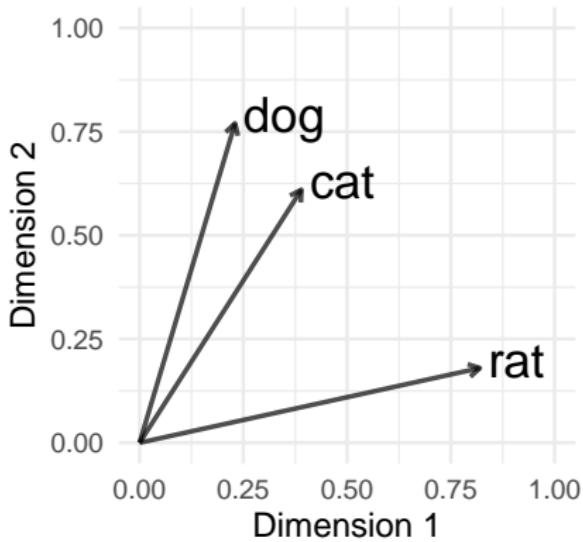
E.g., what if these words can be represented in a two dimensional **embedding** (e.g., two distinct “concepts”)

Then, you might have

- ▷ The word `cat` represented by  $(0.39, 0.61)$
- ▷ The word `dog` represented by  $(0.23, 0.77)$
- ▷ The word `rat` represented by  $(0.82, 0.18)$

Word embeddings have to be *estimated* from a corpus

# Word embeddings



Machine “learns” these dimensions from patterns in corpus

Analyst interprets their meaning

- ▷ Dimension 1 might represent something like “wildness” or “pest-like” traits.
- ▷ Dimension 2 might represent “domesticated companionship.”

You can measure the similarity of words using concepts from linear algebra, like **cosine similarity**

- ▷ Roughly: vectors pointing in same direction are very “similar”
- ▷ One-hot encodings are all perfectly dissimilar using this metric

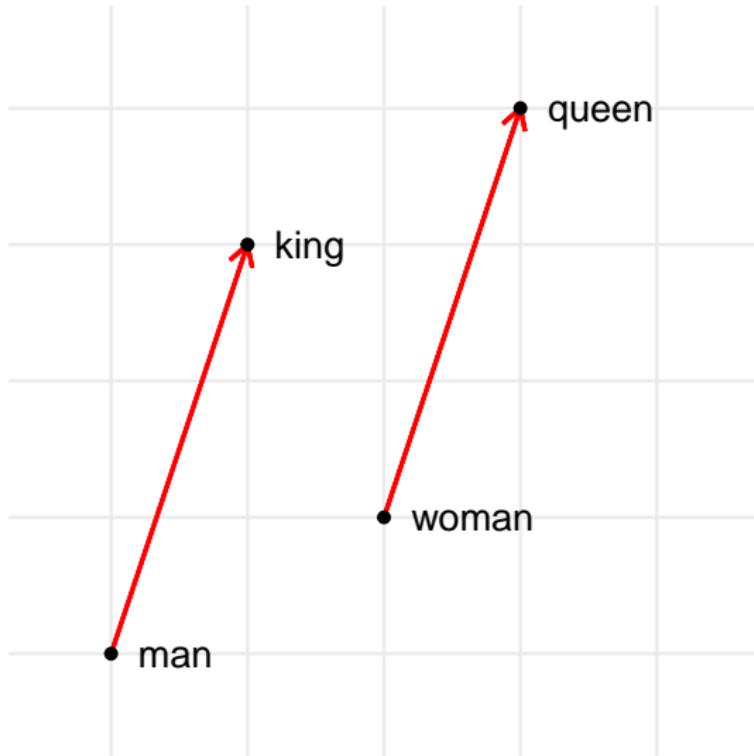
## Word embeddings

Main purpose: use word embeddings as document features instead of words... but why?

1. Technical reason: reduce DFM dimensionality/sparsity
  - ▷ ML models perform better with “denser” vectors
2. Allows for correlations between genuinely related words
  - ▷ One-hot encodings: “dog” and “puppy” are as unrelated as “dog” and “refrigerator”
  - ▷ Embeddings: words like “dog”, “puppy”, “canine” will have similar vectors
  - ▷ Models can generalise better from the data
  - ▷ Capture linguistic relationships through “analogies,”  
like: king - man + woman  $\approx$  queen

# Word embeddings

Analogy: king – man + woman = queen



# Outline

- 1 Text as (quantitative) data
- 2 More complex document features
- 3 Getting text data from other formats
- 4 Data security

# XML

**XML** = eXtensible Markup Language

- ▷ XML has similar structure as HTML, but:
  - XML itself has almost no predefined, domain-specific tags
  - Author can invent XML tags to structure document
- ▷ XML is mostly used to store and distribute data, often as an alternative to JSON (which also has nested structure)
- ▷ HTML is used for *displaying* data (as you know)

Often used to store *textual* data

Reference and further information:

[https://www.w3schools.com/xml/xml\\_whatis.asp](https://www.w3schools.com/xml/xml_whatis.asp)

## XML Example

```
<?xml version="1.0" encoding="UTF-8"?>
<courses>
    <course>
        <title>Data for Data Scientists</title>
        <code>MY472</code>
        <year>2024</year>
        <term>Autumn</term>
        <description>A course about collecting, processing,
                    and storing data.</description>
    </course>
    <course>
        <title>Computer Programming</title>
        <code>MY470</code>
        <year>2024</year>
        <term>Autumn</term>
        <description>An introduction to programming.</description>
    </course>
</courses>
```

## Steps in XML parsing in R

1. Parse an XML file with `read_xml()` in `{xml2}` package
2. Select elements with `html_elements()` or `xml_find_all()`
3. Extract text with `html_text()` or `xml_text()`

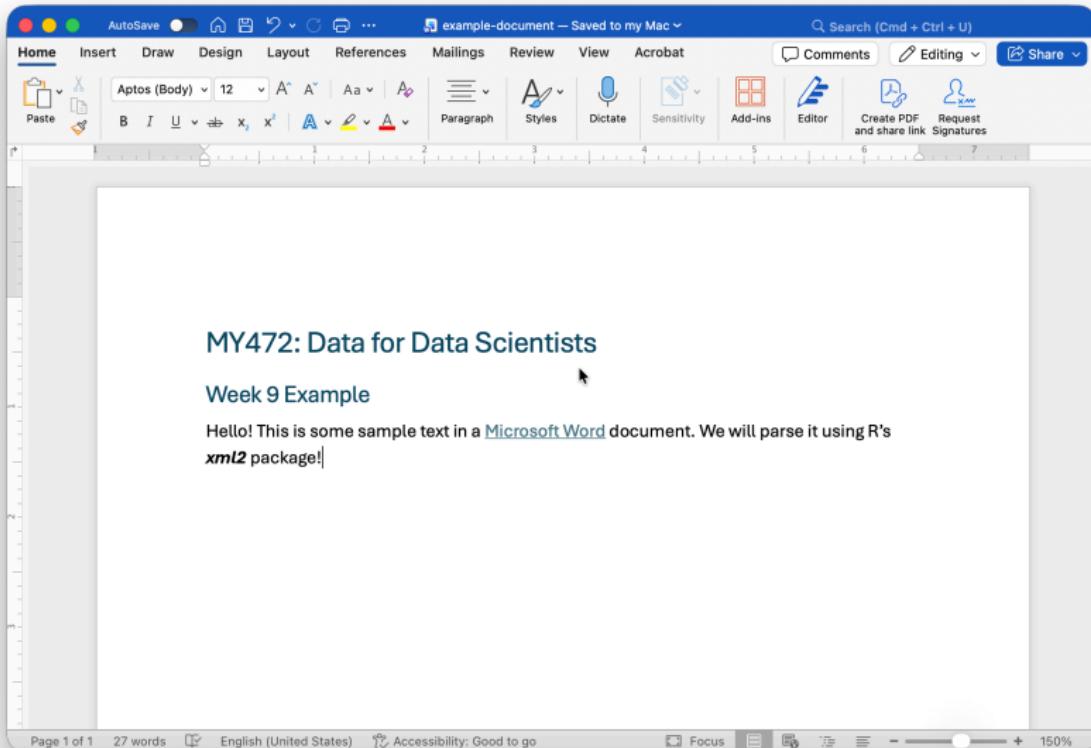
Working with XML in R is nearly identical to working with HTML  
(as we covered in week 7)

- ▷ Just use `{xml2}` to read XML files, and then use `{rvest}` to navigate and parse the tags

## Some XML use cases

- ▷ API data, e.g., Google Maps Geocoding API
- ▷ Data storage, e.g. Canadian members of parliament:  
<https://www.ourcommons.ca/Members/en/search> (select "Export as XML" at the bottom)
- ▷ Scalable Vector Graphics (SVG), e.g. London borough map  
<https://upload.wikimedia.org/wikipedia/commons/b/be/BlankMap-LondonBoroughs.svg>
- ▷ ePub electronic books, e.g. Project Gutenberg books
- ▷ Office documents, e.g. OpenOffice, Microsoft Office
  - Notice difference between .doc and .docx
- ▷ RSS web feeds, e.g. [http://onlinelibrary.wiley.com/rss/journal/10.1111/\(ISSN\)1540-5907](http://onlinelibrary.wiley.com/rss/journal/10.1111/(ISSN)1540-5907)

# XML example: MS Word documents



## XML example: MS Word documents

```
library("rvest")
library("xml2")
word.path <- file.path(ddir, "example-document.docx")
unzip(word.path, exdir = ddir)
unzip(word.path, list = TRUE)
```

		Name	Length	Date
1	[Content_Types].xml		1312	1980-01-01
2	_rels/.rels		590	1980-01-01
3	word/document.xml		3966	1980-01-01
4	word/_rels/document.xml.rels		1003	1980-01-01
5	word/theme/theme1.xml		8391	1980-01-01
6	word/settings.xml		3147	1980-01-01
7	word/styles.xml		43339	1980-01-01
8	word/webSettings.xml		1069	1980-01-01
9	word/fontTable.xml		1831	1980-01-01
10	docProps/core.xml		749	1980-01-01
11	docProps/app.xml		714	1980-01-01

## XML example: MS Word documents

```
word.doc <- read_xml(file.path(ddir, "word/document.xml"))
print(word.doc)

{xml_document}
<document Ignorable="w14 w15 w16se w16cid w16 w16cex w16sdtdh w16sd...
[1] <w:body>\n  <w:p w14:paraId="264D80B7" w14:textId="61" ...
```

## XML example: MS Word documents

```
t <- html_elements(word.doc, xpath = "//w:p")
# t <- xml_find_all(word.doc, xpath = "//w:p") # identical
print(t)

{xml_nodeset (3)}
[1] <w:p w14:paraId="264D80B7" w14:textId="6130D9C8" w:rs ...
[2] <w:p w14:paraId="1B6C64AF" w14:textId="27DAEC61" w:rs ...
[3] <w:p w14:paraId="18DD7CC0" w14:textId="2D8C7BE6" w:rs ...
```

## XML example: MS Word documents

```
t <- html_text(t)
# t <- xml_text(t) # Identical
print(t)
```

```
[1] "MY472: Data for Data Scientists"
[2] "Week 9 Example"
[3] "Hello! This is some sample text in a Microsoft Word document. !"
```

## XML example: MS Word documents

```
t <- t |>  
  paste0(collapse = "\n")  
cat(t)
```

MY472: Data for Data Scientists

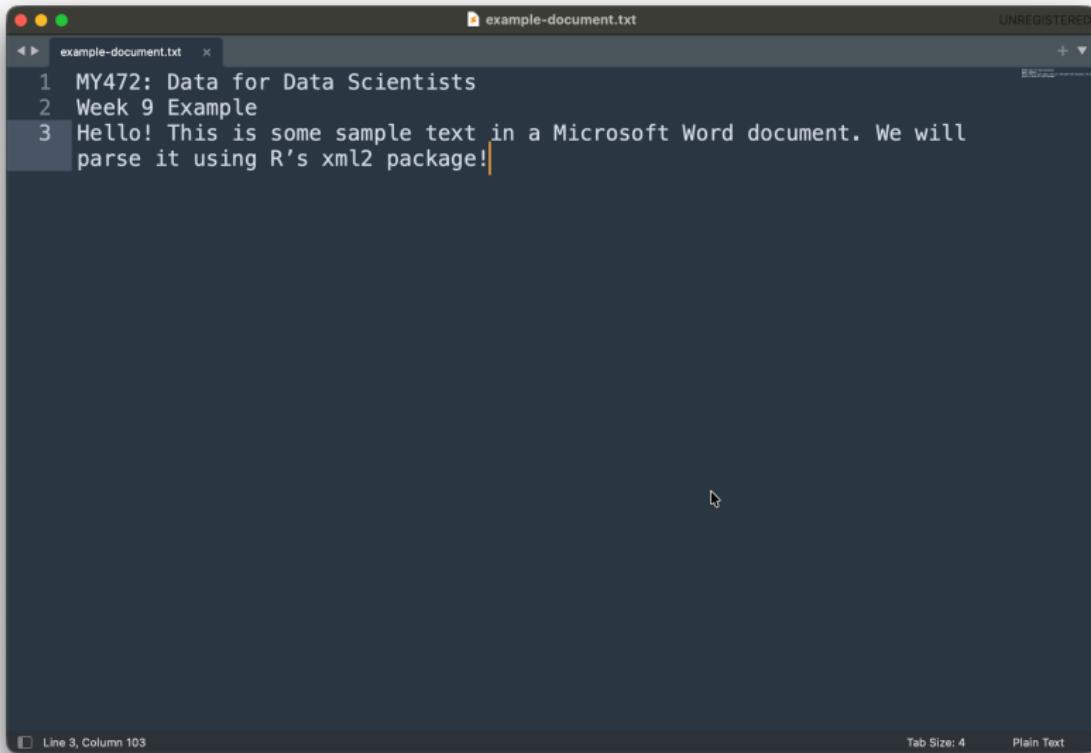
Week 9 Example

Hello! This is some sample text in a Microsoft Word document. We wi

## XML example: MS Word documents

```
library("readr")
library("stringr")
write_file(t, str_replace(word.path, ".docx", ".txt"))
```

# XML example: MS Word documents



A screenshot of a terminal window titled "example-document.txt". The window shows the following text:

```
1 MY472: Data for Data Scientists
2 Week 9 Example
3 Hello! This is some sample text in a Microsoft Word document. We will
   parse it using R's xml2 package!
```

The terminal interface includes standard elements like a title bar, tabs, and status bars at the bottom. The status bar indicates "Line 3, Column 103", "Tab Size: 4", and "Plain Text".

## Other sources of text

Text is available in many other formats, two important ones:

1. **Cloud based documents**: text that is available on a cloud-based editor, e.g., Google Docs
  - ▷ These are usually plain text formats “under the hood”
  - ▷ Often can be exported in some kind of plain text format
  - ▷ There are often APIs to access text directly, such as the [Google Docs API](#)
2. **“Image-like” files**: text is human-readable, might be computer-readable; stores text as images rendered on a page, e.g., [.pdf](#), [.jpg](#), [.png](#), [.tif](#), etc.

# Optical character recognition (OCR)

Image-like files: often how analog texts are digitized

- ▷ For example, scans of old paper documents

PDFs are most common image-like file type for storing text

- ▷ There are different types, see  
<https://pdf.abbyy.com/learning-center/pdf-types/>
- ▷ Some complications around document security

To extract text:

- ▷ PDFs that are digitally generated; can export text directly in e.g., Adobe Acrobat
- ▷ Otherwise, need **optical character recognition (OCR)**
- ▷ Seminar materials: R code for extracting text from PDFs

# Outline

- 1 Text as (quantitative) data
- 2 More complex document features
- 3 Getting text data from other formats
- 4 Data security

## Data security in storage

When you get data from other people, you have an ethical obligation to store it and transfer it securely

- ▷ Especially if it is restricted access data, or it contains personally identifiable information

In academic contexts (e.g., as an LSE student), you *cannot* collect sensitive data without ethics approval

- ▷ LSE-specific information: <https://info.lse.ac.uk/staff/divisions/research-and-innovation/research/Research-integrity-at-LSE>

In UK (and most of Europe), bound by, for example GDPR

Let's look at some of the technical aspects of data security that can help you meet your obligation (but not guarantee it!)

## Data security in storage

Absolute bare minimum: use aggressive authentication practices!

- ▷ Do not disable password/PIN protection
- ▷ Use complex passwords and longer PINs
- ▷ Devices should auto-lock after short period of inactivity
- ▷ Enable biometric authentication (fingerprint, facial recognition)
- ▷ Turn on security settings that protect your digital data if your device is lost or stolen
- ▷ Do not share passwords or user accounts

Two issues:

1. Data security **at rest**: on a storage device
2. Data security **in transit**: transferring over network

## Encryption at rest

Best practice: encrypt digital data

- ▷ **Encryption** is the process by which “meaningful” digital data is transformed into a format that is “unmeaningful”
- ▷ In other words, encryption “locks” your data
- ▷ An **encryption key** is needed to “unlock” (or **decrypt**) encrypted data
- ▷ Core issues: who has the key? is the key secure?

You can encrypt your computer’s hard drive:

- ▷ On macOS, you turn on FileVault
- ▷ On Windows, you turn on BitLocker

(You should do this!)

## Encryption at rest

**Symmetric encryption** is used to encrypt files

- ▷ There is *one* key that encrypts and decrypts
- ▷ If you want multiple people to be able to encrypt/decrypt, need to (securely) share the key

**Asymmetric encryption** is used to share keys

- ▷ Two keys: a public and private key
- ▷ User A's public key + user B's private key → **shared secret**
- ▷ B's public key + A's private key → *same* shared secret
- ▷ More relevant for encryption in transit (more later)

## Encryption at rest

An encryption key is a sequence of randomly generated bytes

- ▷ More bytes → more secure because hackers need to make more guesses
- ▷ Key with 1 byte: can be guessed in less than 256 tries ( $2^8$ )

The key is used for encryption via an encryption algorithm

The modern standard is 256-bit AES, an algorithm that uses keys with 32 random bytes ( $8 \times 32 = 256$  random bits)

- ▷ nearly impossible to break by “brute force”

## Encrypting and decrypting data in R

Can create 256-bit keys in R using `{openssl}` package

```
library("openssl")

# Create a key: sequence of 32 random bytes
key <- rand_bytes(32)
# key <- aes_keygen(length = 32) # equivalent

key
```

```
[1] 2f bb 0b 17 37 c7 7d 01 95 eb 13 31 b7 68 2e 2c 18 de
[19] 33 a4 14 f5 86 a5 29 16 84 49 c0 77 88 b0
```

# Encrypting and decrypting data in R

If you plan to use this as an encryption key, you better not lose it!

This is a very insecure way to store it:

```
writeBin(key, "my_key.bin")
```

Saving to keychain is much, much better:

```
# Save key to keychain
library("keyring")
key_set_with_raw_value("test_enc_key", password = key)
# Delete key
rm(key)
```

But be careful: if someone has access to R on your computer, they may be able to get this key!

- ▷ In seminar, you'll learn how to lock keychains

# Encrypting and decrypting data in R

To encrypt text using R (and a key), you first convert the text to raw bytes, then encrypt

```
# What's my secret message
secret <- "A top secret message!"

# Convert to bytes
secret <- charToRaw(secret)

# Encrypt the message using key
secret <- aes_gcm_encrypt(secret,
                           key = key_get_raw("test_enc_key"))
rawToChar(secret) # Secret has been encrypted!
```

```
[1] "\xe6\xbd\"022V\xe9\xb6\003\xe1\035,\xbbHiy\002)I"
```

Note: you can encrypt *any* digital data, not just text

# Encrypting and decrypting data in R

To decrypt the message, you need the key

```
# Decrypt the message using key (pulled from keychain)
secret <- aes_gcm_decrypt(secret,
                           key = key_get_raw("test_enc_key"))
rawToChar(secret)
```

```
[1] "A top secret message!"
```

## Encryption in transit

Lots of your digital data lives off your device

- ▷ E.g., cloud storage, emails, text messages, photos, etc.
- ▷ This is becoming even more common, e.g. “optimised” storage on Apple devices and default sync settings in Dropbox

You can control security of your local device, but not remote devices (the “cloud”) or network protocols

Bare minimum requirements for cloud storage:

- ▷ Encryption in transit using **Transport Layer Security (TLS)**
- ▷ Encryption at rest (on server) using AES 256-bit encryption

Reputable providers do this (often for legal reasons)

## Encryption in transit

How does this apply to web browsing?

- ▷ `http` protocol does not use encryption in transit (yikes!), but `https` protocol does
- ▷ Do not put your sensitive information (like credentials) into a webpage with an `http://` url
- ▷ If you have a website, make sure it uses `https` or many browsers may block access to it

Private browsing does *not* protect your data in transit

- ▷ But it can be useful for other things: avoiding web tracking, limiting browsing history, etc.

## Encryption in transit

Gold standard is **end-to-end (E2E) encryption**:

- ▷ Encrypt data on local device *before* transit
- ▷ Decrypt on remote device *after* transit
- ▷ Core challenge: how to share encryption key between sender and recipient

Examples: Signal, WhatsApp, many **iCloud** services, Apple's **Advanced Data Protection**

Most cloud storage services do not use/offer E2E encryption

- ▷ Dropbox does not use E2EE!
- ▷ Only major provider offering E2EE to individual consumers is Apple iCloud (but **not in the UK...yikes!**)

## Encryption in transit

Still: hackers are always working to break encryption

- ▷ For example, [recent reports](#) of a trojan that can read decrypted Signal and WhatsApp messages on Android devices

Best practices:

- ▷ Regularly delete old messages and old files you no longer need
- ▷ Turn on auto-delete for your messaging apps (if available)
- ▷ Keep up to date on the security of the services you use
- ▷ Buy devices and use operating systems and software that are known to be more secure

## Encryption in transit

You can set up a simple E2E encryption pipeline in R

You will see an example in seminar, but basic structure:

- ▷ Person A and Person B create **asymmetric keys** with public and private components
- ▷ Share the public keys with one another
- ▷ When A creates an encrypted message:
  - Use B's public key + A's private key to make a **shared secret**
  - Use shared secret to create a symmetric encryption key that is then used to encrypt the message
  - Send the encrypted message
- ▷ Person B uses their private key + A's public key to decrypt the symmetric key and decrypt the message

## Data security: other considerations

So far, narrowly focused on technical aspects of digital data security assuming otherwise benevolent cloud providers

But companies might be accessing your data

- ▷ LLM providers like OpenAI are opaque about whether/how they use personal data you give them
- ▷ Gmail uses email data to personalise services
- ▷ Cloud storage providers have access to the information in your private files (except iCloud, if ADP is turned on)

Companies are sometimes vague or misleading

- ▷ E.g., WhatsApp chat backups do not use E2EE by default

To be fair: consumer demands are partly to blame

# Encryption can be controversial

Encryption can be controversial

Some governments do not want individuals to strongly encrypt their digital data (e.g., the UK government and Apple devices)

So, traveling with encrypted data poses a risk

- ▷ You may be required to decrypt your digital data at borders
- ▷ Laws vary from jurisdiction to jurisdiction
- ▷ Your ethical obligations are not waived simply because you did not do your due diligence before transporting data

LSE's guidelines: <https://info.lse.ac.uk/staff/divisions/dts/assets/documents/policies/Guidelines-Encryption-Guidelines-v1-1.pdf>

## What should you do?

What's the take-away for you as a data scientist?

- ▷ Start by shoring up your digital security
- ▷ To be safest: never store sensitive data off your (fully encrypted) devices
- ▷ If you still need to use cloud services:
  - Use services provided by your institution, that have gone through security and legal reviews
  - For E2E encryption, use iCloud (with ADP enabled) or manually encrypt or use third-party software
- ▷ Spend some time reviewing policies
- ▷ Learn security practices of your service providers